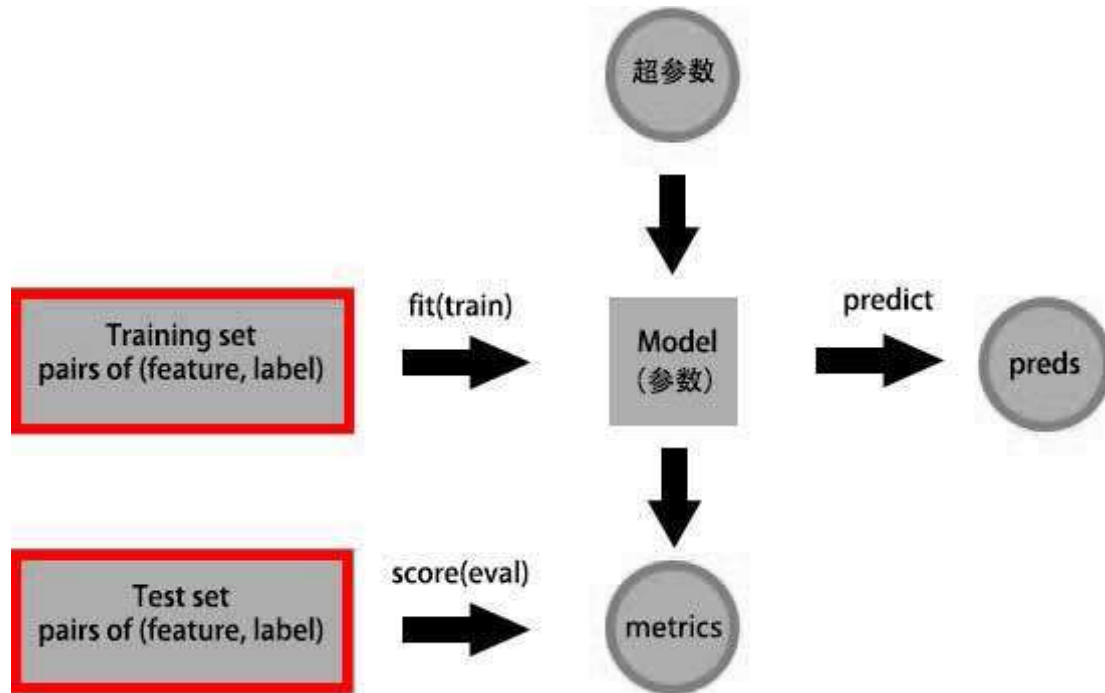


In [10]:

```
%matplotlib inline
```

feature engineering

寻找描述数据最恰当的feature



1. 数据的标准化（连续型参数）

最基本的一种预处理手段。各特征若scale差异很大，可能在训练的数值运算过程中产生显著bias，standardization的目的是去除这种由于scale显著差异而导致的bias

通过线性变换（平移和rescale），使各维度的参数中心为0，标准差为1。可视为一种 trivial unsupervised learning，在 sklearn 中用 transformer表示。

tip：当feature各维度是连续的，含义差别大，尺度差别大，且没有更多思路时，总可以先进行standardization，standardization 确保各feature的scale大体一致。如同，我们对事件发生的概率没有更多信息时，总可以先假定事件发生概率是均匀分布的。

In [15]:

```
from sklearn.preprocessing import scale
import numpy as np

data = np.array([
[0.06, 1.0, 52, 1050],
[0.14, 1.4, 214, 3481],
[0.02, 2.6, 78, 4229]
])

x = scale(data) # mu = 0, var = 1

# print("x=%s" % x)

# print("data mean=%s" % np.mean(data, axis=0))
# print("data var=%s" % np.var(data, axis=0))

# print("x mean=%s" % np.mean(x, axis=0))
# print("x var=%s" % np.var(x, axis=0))
```

In [6]:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

iris = load_iris()
X, y = iris.data, iris.target

train_X, test_X, train_y, test_y = train_test_split(X, y, train_size=0.8, stratify=y)

mean_X = train_X.mean(axis=0)
std_X = train_X.std(axis=0)

# 手动standardization
train_X_standardized = (train_X - mean_X) / std_X
train_X_standardized.mean(axis=0)
train_X_standardized.std(axis=0)

# 自动scaler
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(train_X)
train_X_standardized = scaler.transform(train_X)
test_X_standardized = scaler.transform(test_X)
```

注意：有的时候同类型 feature“自然的”variance差异，应当保留

2. dimensionality reduction (数据降维)

- [PCA \(Principal Component Analysis\) \(PCA.ipynb\)](#)
- 其他降维手段：feature selection (寻找和目标)

In [1]:

```

from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import numpy as np

def plot_pca_illustration():
    rnd = np.random.RandomState(5)
    X_ = rnd.normal(size=(300, 2))
    X_blob = np.dot(X_, rnd.normal(size=(2, 2))) + rnd.normal(size=2)

    pca = PCA()
    pca.fit(X_blob)
    X_pca = pca.transform(X_blob)

    S = X_pca.std(axis=0)

    fig, axes = plt.subplots(2, 2, figsize=(10, 10))
    axes = axes.ravel()

    axes[0].set_title("Original data")
    axes[0].scatter(X_blob[:, 0], X_blob[:, 1], c=X_pca[:, 0], linewidths=0,
                    s=60, cmap='viridis')
    axes[0].set_xlabel("feature 1")
    axes[0].set_ylabel("feature 2")
    axes[0].arrow(pca.mean_[0], pca.mean_[1], S[0] * pca.components_[0, 0],
                  S[0] * pca.components_[0, 1], width=.1, head_width=.3,
                  color='k')
    axes[0].arrow(pca.mean_[0], pca.mean_[1], S[1] * pca.components_[1, 0],
                  S[1] * pca.components_[1, 1], width=.1, head_width=.3,
                  color='k')
    axes[0].text(-1.5, -.5, "Component 2", size=14)
    axes[0].text(-4, -4, "Component 1", size=14)
    axes[0].set_aspect('equal')

    axes[1].set_title("Transformed data")
    axes[1].scatter(X_pca[:, 0], X_pca[:, 1], c=X_pca[:, 0], linewidths=0,
                    s=60, cmap='viridis')
    axes[1].set_xlabel("First principal component")
    axes[1].set_ylabel("Second principal component")
    axes[1].set_aspect('equal')
    axes[1].set_ylim(-8, 8)

    pca = PCA(n_components=1)
    pca.fit(X_blob)
    X_inverse = pca.inverse_transform(pca.transform(X_blob))

    axes[2].set_title("Transformed data w/ second component dropped")
    axes[2].scatter(X_pca[:, 0], np.zeros(X_pca.shape[0]), c=X_pca[:, 0],
                    linewidths=0, s=60, cmap='viridis')
    axes[2].set_xlabel("First principal component")
    axes[2].set_aspect('equal')
    axes[2].set_ylim(-8, 8)

    axes[3].set_title("Back-rotation using only first component")
    axes[3].scatter(X_inverse[:, 0], X_inverse[:, 1], c=X_pca[:, 0],
                    linewidths=0, s=60, cmap='viridis')
    axes[3].set_xlabel("feature 1")
    axes[3].set_ylabel("feature 2")
    axes[3].set_aspect('equal')

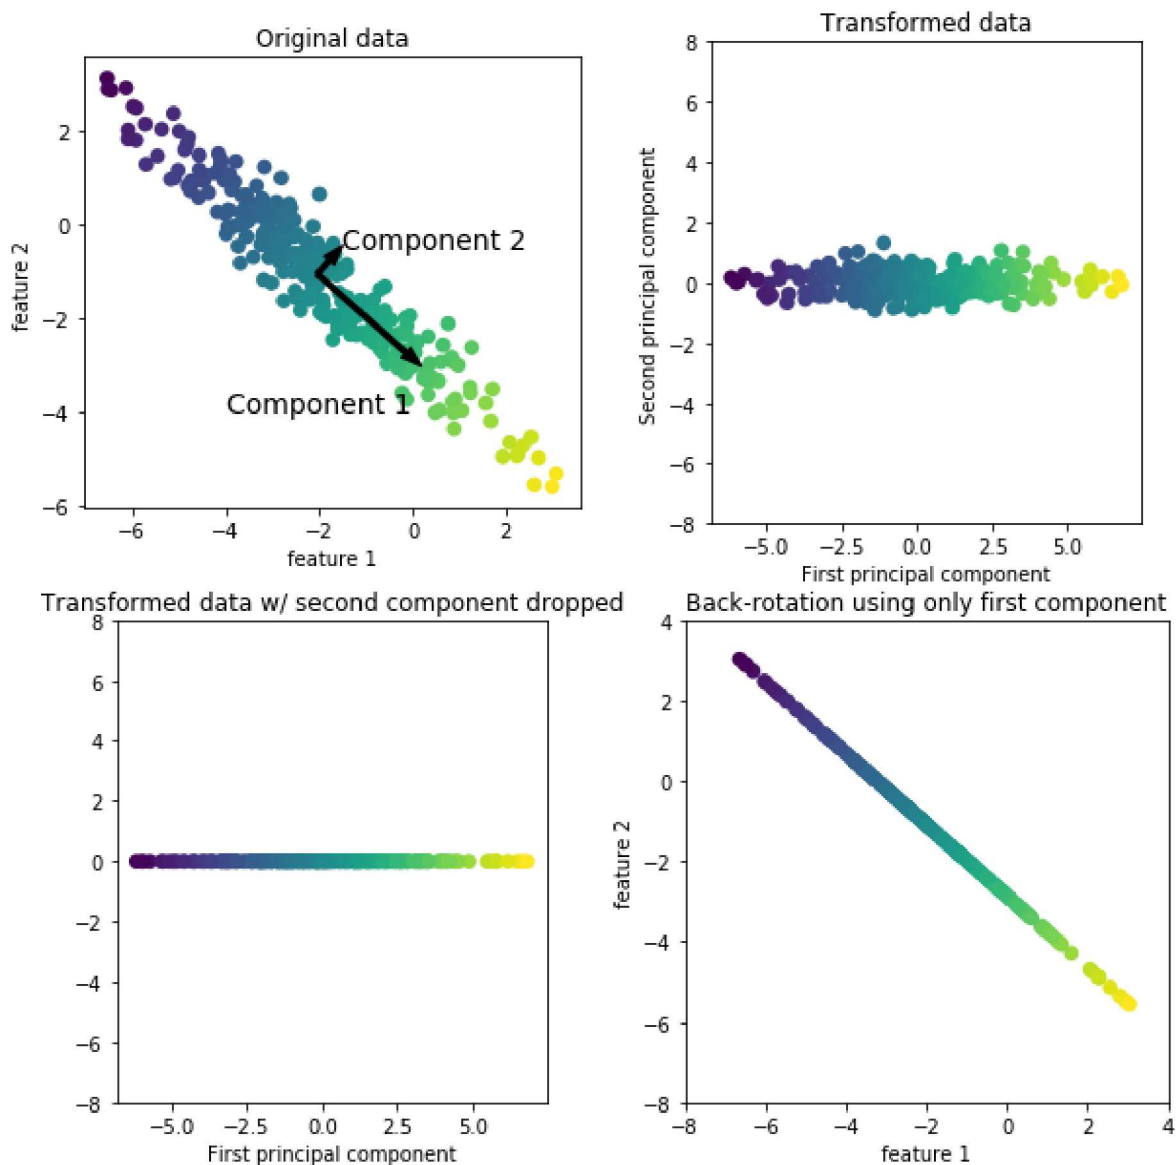
```

```

axes[3].set_xlim(-8, 4)
axes[3].set_ylim(-8, 4)

plot_pca_illustration()

```



2. feature representation

对于简单的numerical参数, representation是其自身

2.1 categorical feature: one-hot 编码

- ordinal: 如尺码、评分等级 (discreted numerical, 假冒的类别型特征)
- nominal: 如名称、性别

why not 0, 1, 2?

携带了不必要的大小排序信息, 彼此关系不等价

In [1]:

```

measurements = [
    {'drink': 'water'},
    {'drink': 'milk'},
    {'drink': 'juice'},
]

# measurements = [
#     {'drink': 'water', 'food': 'beef'},
#     {'drink': 'milk', 'food': 'cake'},
#     {'drink': 'juice', 'food': 'pizza'},
# ]

from sklearn.feature_extraction import DictVectorizer

vec = DictVectorizer()

vec.fit_transform(measurements).toarray() # 稀疏矩阵转常规矩阵

```

Out[1]:

```

array([[ 0.,  0.,  1.],
       [ 0.,  1.,  0.],
       [ 1.,  0.,  0.]])

```

In [19]:

```
vec.get_feature_names()
```

可能遇到的问题：

- 维度空间太高，数据太稀疏，对策：稀疏矩阵
- 不能反映样本点之间的潜在亲疏关系，对策：embedding

问题：既有数值型参量，又有类别性参量怎么办

2.2 文本：bag-of-words encoding（词袋模型）

一些基本概念

corpus（语料）：document（文案，如一段话或一篇文章）的集合 D ，文案数记为 N vocabulary（词库）：number of unique words，词库条目数记为 K

词袋模型：每个document用 K 维向量描述

denote the "bare-frequency" vector of document d as x (the value for term t is denoted by $f(t, d)$)

term frequency vector of x' of document d for word t is given by, $tf(t, d) \equiv \frac{f(t, d) + 1}{|x|}$, 1 is for smoothing, denominator for normalizing long document.

some variant, such as, "logarithmically scaled term frequencies" $tf(t, d) \equiv \log(f(t, d) + 1)$, to mitigate long document

inverse document frequency $idf(t, D) \equiv \log \frac{N}{1+n}$, N is the size of the corpus D , n is the number documents which contain term t .

表征一句话或一个document的步骤:

1. 预处理: (中文需要先分词) 生僻词用UNK代替, stop words过滤 (英语the a and, 中文了, 的, 着等), 同义词合并 (text to lowercase, 同根词合并, 如running run) (起到 dimensionality reduction 效果)
2. 通过计算某个document中所有词汇的tf-idf 值, 给出表征该document的向量

hashing tricks, 节省下词汇dict的存储空间, 同时不比提前preview所有document, 因而使用于“在线”学习问题 (cons: 可能存在hashing 碰撞)

In [3]:

```
X = ["Some say the world will end in fire,",
     "Some say in ice."]

from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()

vectorizer.fit(X)
print(vectorizer.vocabulary_)
print(vectorizer.get_feature_names())
```

```
{u'end': 0, u'fire': 1, u'some': 5, u'ice': 2, u'will': 7, u'say': 4, u'in': 3, u'wo
rld': 8, u'the': 6}
[u'end', u'fire', u'ice', u'in', u'say', u'some', u'the', u'will', u'world']
```

In [4]:

```
X_bag_of_words = vectorizer.transform(X) # sparse matrix
print(X_bag_of_words.toarray())
vectorizer.inverse_transform(X_bag_of_words)
```

```
[[1 1 0 1 1 1 1 1 1]
 [0 0 1 1 1 1 0 0 0]]
```

Out[4]:

```
[array([u'end', u'fire', u'in', u'say', u'some', u'the', u'will', u'world'],
      dtype='<U5'), array([u'ice', u'in', u'say', u'some'],
      dtype='<U5')]
```

In [5]:

```

from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer()

vectorizer.fit(X)
print(vectorizer.vocabulary_)
print(vectorizer.get_feature_names())
X_bag_of_words = vectorizer.transform(X) # sparse matrix
print(X_bag_of_words.toarray())
vectorizer.inverse_transform(X_bag_of_words)

```

注意到词语排序信息丢失的问题

```

{u'end': 0, u'fire': 1, u'some': 5, u'ice': 2, u'will': 7, u'say': 4, u'in': 3, u'wo
rld': 8, u'the': 6}
[u'end', u'fire', u'ice', u'in', u'say', u'some', u'the', u'will', u'world']
[[ 0.39166832  0.39166832  0.          0.27867523  0.27867523  0.27867523
    0.39166832  0.39166832  0.39166832]
 [ 0.          0.          0.63009934  0.44832087  0.44832087  0.44832087
    0.          0.          0.          ]]

```

Out [5]:

```

[array([u'world', u'will', u'the', u'some', u'say', u'in', u'fire', u'end'],
      dtype='<U5'), array([u'some', u'say', u'in', u'ice'],
      dtype='<U5')]

```

可能遇到的问题:

- 中文有分词问题 (编码问题, 对策 #encoding=utf-8)
- 忽略了次序信息, 对策: 采用sequence描述一句话, 序列中的每个词用one-hot编码或者用embedding 编码

In [16]:

```
#encoding=utf-8

import jieba
import jieba.posseg as pseg
import numpy as np
import pandas as pd

df = pd.read_csv('../data/spams.csv', delimiter='\t+', header=None)

corpus_raw = df[1]
y = df[0]

def preprocess_doc(doc):
    words=list(pseg.cut(doc))
    words_filtered = filter(lambda w: w.flag!= "x" and w.flag!="m" and (w.word not in ["请", "白
    return " ".join(map(lambda w: w.word, words_filtered))

corpus = map(preprocess_doc, corpus_raw)

for d in corpus:
    print(d)
```

C:\Users\Lenovo\Anaconda2\lib\site-packages\ipykernel_launcher.py:8: ParserWarning:
Falling back to the 'python' engine because the 'c' engine does not support regex se
parators (separators > 1 char and different from '\s+' are interpreted as regex); yo
u can avoid this warning by specifying engine='python'.

Building prefix dict from the default dictionary ...
Dumping model to file cache c:\users\lenovo\appdata\local\temp\jieba.cache
Loading model cost 2.721 seconds.
Prefix dict has been built succesfully.

您好 代办 各类 证件 文凭 信用卡 咨询
无需 担保 就可 申请 高额 透支 信用卡 联系 王先生
代办 各类 证件 催账 联系 周先生
尊敬 用户 您 流量 包已 到期 详情 咨询
老板 说 要 今天 要 赶 下工 不 回去 吃饭 哈
今天 去 哪 吃饭 庆祝
代办 信用卡 高利贷 无需 担保 联系 李小姐
话费 快用 完 帮 交

2.3 图片:

- 纯像素grayscale/rgb 值向量
- sift提供矢量 (Scale-Invariant Feature Transform) 或 Speeded-Up Robust Features (SURF) descriptor
- CNN 自动提取特征

In [12]:

```
from skimage.io import imread, imshow, show

data = imread("images/neo.png")

print(data)

imshow(data)
show()
```